From Invariant to Implementation: Enforcing Neurorights and Eco-Governance in a Regulated Bioscale Safety System

Architectural Foundations: The Dual-Track Model of Contracts and Toolchain

The bioscale safety system is architecturally defined by a dual-track model that tightly couples governance through machine-readable contracts with technical enforcement via a robust toolchain . This design philosophy moves beyond traditional policy documents to create an executable specification, ensuring that high-level ethical and safety obligations are translated into low-level, verifiable software constraints. The entire system is anchored in a set of core invariants encoded once in Application Language Notation (ALN) and enforced everywhere, creating a "focus safety constitution" that governs all components . This foundational layer consists of three primary ALN shard entities: rights contracts, safety and eco profile shards, and a central regulatory profile shard that ties them together .

The first track, governed by ALN contracts, establishes the system's operational boundaries and ethical commitments. These contracts are not merely data schemas but serve as the system's constitutional framework, defining what is permissible and what must be prevented . Key entities within this track include `focus.rights.contract.v1` and `focus.safety.eco.profile.v1` . The `focus.rights.contract.v1` shard encodes critical augmented rights invariants, such as the principle that only bounded, dimensionless metrics may leave the device, while raw EEG or behavioral streams remain confined to local storage . It also enforces that every adaptation, whether it be UI simplification or a break prompt, must be optional and reversible, ensuring the system can be disabled to restore baseline functionality without degradation . This directly implements principles of cognitive liberty and mental integrity. The `focus.safety.eco.profile.v1` shard, conversely, focuses on physical and environmental safety, aggregating daily focus minutes, breaks taken, and device-hours reduced into metrics like the `EcoImpactScore` . This shard becomes the canonical evidence that the system operates within healthy and eco-friendly envelopes . The interplay between these two shards ensures that no action is taken that would violate either the user's rights or the system's safety and environmental goals. For instance, a decision to suggest a break is validated against both the fatigue index (a rights/safety metric) and the cumulative impact on device hours (an eco-metric).

The second track, the toolchain, provides the mechanism to enforce these contractual invariants. Its purpose is to make unsafe behavior literally unrepresentable in safe code, thereby grounding the abstract promises of the contracts in concrete, provably correct software . This is achieved primarily through the use of Rust and C++, leveraging their strong type systems and memory safety guarantees to build a defense-in-depth architecture . The Rust-based `bioscale-metrics` crate is central to this effort, generating bounded types like `FocusIndex` and `FatigueIndex` that enforce the `[0,1]` range at construction . Any attempt to create an instance of these types with an out-of-bounds value would fail at compile time, making invalid data an impossibility in the system's safe execution path . Similarly, the C/C++ backend modules are designed to perform one-way aggregation, transforming raw neurodata into the bounded metrics specified in the ALN shards and discarding reconstruction details to reinforce privacy . This dual-track approach ensures that the system's governance is not just declared in documentation but is embedded within its very structure, providing a solid foundation upon which human trials and audits can be built . The `focus.regulatory.profile.v1` shard serves as the nexus of this architecture, acting as a central, machine-readable "regulatory label" that explicitly maps each module to its governing rights, safety profiles, and relevant regulatory references, including those from the EU AI Act and neurorights frameworks .

ALN Shard	Purpose	Key Invariants & Metrics
`focus.rights.contract.v1`	Defines user rights, data handling policies, and opt-out mechanisms.	Raw data never leaves device; only bounded aggregates exit. All adaptations are optional and reversible. Local processing is mandatory.
`focus.safety.eco.profile.v1`	Tracks safety and environmental impact metrics over time.	Aggregates daily focus minutes, breaks, and device-hours reduced. Calculates `AvgDailyDeviceHoursReduced` and `EcoImpactScore`.
`focus.regulatory.profile.v1`	Acts as a central hub linking modules to their governance context.	Maps module IDs to rights contracts and safety profiles. References EU AI Act articles and neurorights principles.

This architectural model represents a sophisticated synthesis of software engineering and governance theory. By encoding invariants directly into the system's data structures and enforcement mechanisms, it creates a system that is, by design, aligned with its intended ethical and safety goals. The tight coupling between the ALN contracts and the Rust/C+

+ toolchain ensures that any deviation from these contracts is a compile-time error, preventing regressions and ensuring long-term stability and trustworthiness.

Technical Enforcement: Making Contract Violations Unrepresentable in Rust/C++

The technical enforcement layer of the bioscale safety system is engineered to transform the abstract legal and ethical invariants defined in ALN contracts into concrete, unbreakable software guarantees. This is accomplished through a carefully selected stack of technologies and programming paradigms, primarily centered on Rust and C++, that leverage strong static typing and memory safety to make contract violations impossible to represent in memory . The strategy is not to check for violations at runtime, but to prevent their creation in the first place, achieving a higher degree of safety than is possible with traditional testing alone.

A cornerstone of this enforcement strategy is the development of a dedicated Rust crate, tentatively named `qpudatashards/libbioscale_metrics` . This crate's primary function is to generate a suite of bounded types directly from the ALN schema definitions. For example, the ALN-defined constraint that `FocusIndex` and `FatigueIndex` must lie within the `[0,1]` interval is implemented as a compile-time guarantee within the `FocusIndex` and `FatigueIndex` structs in Rust . Construction of these types would involve a validation step; any value outside the specified bounds would cause the program to fail to compile. This technique effectively eliminates entire classes of bugs related to out-of-bounds data, such as incorrect thresholds for break prompts or invalid calculations for the `EcoImpactScore`. Once constructed, these types are used consistently across all guards, loggers, and business logic, ensuring that downstream components operate exclusively on validated, well-formed data . This prevents corruption or invalidation of metrics throughout the system's lifecycle. Furthermore, this approach allows for the normalization of disparate metrics—for instance, converting raw kilowatt-hours and device hours into a single, dimensionless `EcoImpactScore`—which can then be compared and aggregated consistently, enabling cross-study comparisons and standardized eco-impact accounting .

Complementing the Rust-based safe codebase is a C/C++ backend responsible for the initial transformation of raw data . This backend ingests non-invasive signals, such as consumer-grade EEG, sEMG, keystroke dynamics, and app usage patterns, and performs the necessary computations to derive the bounded metrics defined in the ALN shards .

The design of this component is critical for maintaining the system's privacy guarantees. It is engineered for one-way aggregation, meaning that after a bounded metric is calculated, the detailed information required to reconstruct the original raw signal is discarded . This irreversible transformation reinforces the invariant that raw neural data never leaves the local device, a principle essential for protecting mental privacy . Unit tests and schema validators are integrated into the CI pipeline to ensure that any attempt to write raw neurodata or user identifiers to ALN logs fails immediately, providing an additional layer of defense against accidental or malicious data leakage . The combination of C/C++ for secure data processing and Rust for subsequent, provably safe computation creates a robust and layered security architecture.

To provide empirical validation of the system's dynamic behavior, a specialized property-testing harness named `bioscaletest` is developed . While the bounded types provide formal, compile-time guarantees, property tests offer a way to verify that the system's runtime behavior adheres to the intended invariants under a wide variety of conditions. This harness simulates long sessions and stress scenarios to check for emergent properties. For instance, one key property test verifies that break prompts always trigger before the `FatigueIndex` exceeds a critical threshold (e.g., 0.9) or before the session length surpasses the cap defined in the rights contract . Another test checks for monotone eco-benefits by ensuring that the `EcoImpactScore` never decreases when the cohort's `AvgDailyDeviceHoursReduced` improves for the same period . Finally, the harness includes tests to confirm that no code path attempts to act when the rights contract explicitly flags a violation of local processing or explainability, serving as a final sanity check on the system's adherence to its core principles . Together, the bounded types, the secure C/C++ backend, and the property-testing harness form a comprehensive technical enforcement strategy that makes contract violations both statically and dynamically impossible, building a foundation of provable safety and rights-respecting operation.

Governance Alignment: Mapping Technical Artifacts to Global Regulatory Frameworks

The bioscale safety system is not only technically sound but is also architected to align with several major global governance frameworks, including the EU AI Act, emerging neurorights legislation exemplified by Chile, and the recommendations of the OECD and UNESCO. This alignment is not an afterthought but is woven into the fabric of the system's design, where each technical artifact has a clear justification in terms of its contribution to regulatory compliance and ethical responsibility. The `focus.regulatory.profile.v1` ALN shard serves as a master map, explicitly linking modules to their corresponding regulatory obligations .

The system's design demonstrates a remarkable congruence with the EU AI Act, particularly its provisions for high-risk AI systems [12] [64] . Article 9 mandates a continuous risk management system throughout the AI lifecycle; the entire bioscale architecture—from the ALN-defined risk profiles to the CI validators that enforce monotone safety—is a direct implementation of this requirement [37] . Monotone safety, which constrains OTA updates so that safety bounds cannot be loosened (`B_{new} ≤ B_{old}`), is a concrete technical mechanism for managing risk over time . Article 10 requires robust data governance and the prevention of discriminatory outcomes. The system's core invariant of processing only bounded aggregates locally and never sharing raw data aligns with the need for representative datasets and strong data protection [35] . The planned fairness analysis, which will examine break rates and throttle decisions across demographic groups using de-identified ALN statistics, is a direct response to the Act's mandate to mitigate bias [35] . Furthermore, Article 14's requirement for meaningful human oversight is satisfied by the explicit guarantee that all system adaptations are optional and reversible, allowing users to maintain ultimate control [35] . The focus on tracking environmental metrics like `AvgDailyDeviceHoursReduced` and `EcoImpactScore` directly addresses Article 27's obligation to consider societal and environmental well-being [35] .

In its approach to neurorights, the system anticipates and meets the standards set by pioneering jurisdictions like Chile. The 2021 constitutional amendment in Chile established new rights to protect mental privacy, cognitive liberty, and mental integrity against adverse effects from neurotechnologies [2] . The bioscale system's invariant that raw EEG and behavioral data never leave the local device is a direct technological implementation of the right to mental privacy, preventing the unauthorized access or profiling of cerebral activity [2] [4] . The requirement that all adaptations be reversible upholds cognitive liberty, ensuring users can always opt-out and restore their baseline experience without permanent consequences [2] [20] . The emphasis on non-coercive interventions and the ability to completely disable the system safeguards mental integrity, protecting users from manipulation and harm to their personal identity [2] [21] . This proactive design positions the system to comply with Chile's strict liability regime, where producers and administrators are held jointly and strictly liable for any damage caused by their products [2] .

Finally, the project's ethos is deeply aligned with the guiding principles articulated by international organizations. The OECD Recommendation on Responsible Innovation in Neurotechnology calls for anticipatory governance and the management of risks associated with technology convergence [16] [50]. The bioscale system's modular, contract-driven architecture embodies this by allowing for incremental evolution while anchoring each change to a pre-defined safety envelope (`bio.safety.envelope.citizen.v1`). The emphasis on transparency through audit-ready logs and the focus on responsible development reflect the OECD's broader guidance on fostering scientific collaboration and public trust [17] [55]. Similarly, the system strongly resonates with the UNESCO Recommendation on the Ethics of Neurotechnology, which aims to safeguard the human mind and establish clear ethical boundaries [19] [46]. The system's core design, which treats brain-adjacent data as highly sensitive and implements strong privacy-by-design features, is a direct response to UNESCO's warnings about threats to mental privacy and autonomy from neurotechnological applications [22] [46].

Regulatory Framework	Principle / Requirement	Technical Artifact / Mechanism	ALN Shard Reference
EU AI Act	Article 9: Lifecycle Risk Management	Monotone safety invariants (`B_new ≤ B_old`) enforced in CI.	`focus.safety.eco.profile.v1` [37]
EU AI Act	Article 10: Data Governance & Non-Discrimination	Processing only bounded aggregates locally; fairness analysis on de-identified logs.	`focus.rights.contract.v1` [35]
EU AI Act	Article 14: Human Oversight	All adaptations are optional and reversible; system can be fully disabled.	`focus.rights.contract.v1` [35]
EU AI Act	Article 15: Robustness & Safety	Property tests (`bioscaletest`) simulate stress, drift, and adversarial conditions.	N/A [35]
EU AI Act	Article 27: Societal Well-being	Tracking and reporting of `AvgDailyDeviceHoursReduced` and `EcoImpactScore`.	`focus.safety.eco.profile.v1` [35]
Chile Neurorights	Mental Privacy	Raw data never leaves device; only bounded aggregates are processed/shared.	`focus.rights.contract.v1` [2] [45]
Chile Neurorights	Cognitive Liberty	All system adaptations are optional and reversible by the user.	`focus.rights.contract.v1` [2] [20]
Chile Neurorights	Mental Integrity	Interventions are non-coercive; disabling the system causes no degradation.	`focus.rights.contract.v1` [2] [21]
OECD Neurotech Rec.	Anticipatory Governance	Modular design with `cyberlinks` anchors to a `bio.safety.envelope.citizen.v1`.	`focus.regulatory.profile.v1` [50] [61]
UNESCO Neurotech Rec.	Safeguarding the Mind	Core design treats brain-adjacent data as highly protected sensitive information.	`focus.rights.contract.v1` [19] [46]

This deliberate mapping demonstrates that the bioscale safety system is more than a piece of technology; it is a governance artifact designed to be compliant by construction.

Operational Validation: Human Trials, Audits, and Continuous Monitoring

The theoretical soundness and regulatory alignment of the bioscale safety system must ultimately be validated through empirical evidence gathered from real-world use, continuous monitoring, and rigorous auditing. The research plan outlines a multi-layered validation strategy that progresses from controlled pilot studies to large-scale, auditable deployments, ensuring that the system performs as intended and continues to uphold its safety and rights-preserving invariants over time . This operational phase is critical for building trust and demonstrating the system's efficacy and safety to regulators, users, and independent reviewers.

The initial stage of validation involves a controlled human study, specifically a Phoenix eco-focus pilot . This study is designed to recruit a small cohort of volunteers (50–200 participants) over a defined period (4–8 weeks) to measure the system's impact in a real-world setting . The primary endpoints of this pilot are the reduction in average daily device hours (`AvgDailyDeviceHoursReduced`) and self-reported strain, with a secondary goal of demonstrating non-inferior or improved performance on attention-related tasks . These metrics directly correspond to the `EcoImpactScore` and health engagement indices tracked in the `focus.safety.eco.profile.v1` shard, providing tangible evidence of the system's benefits . The data collected from this pilot will be logged in ALN format, preserving a complete record of the system's state, decisions, and the resulting outcomes, which will be instrumental for subsequent analysis and refinement .

Parallel to the human trials, a series of neurorights audit exercises are planned to assess the system's transparency and interpretability . In these exercises, independent reviewers will be provided with de-identified ALN logs containing only the bounded metrics, the active contracts, and the system's decisions (e.g., why a break was prompted) . The reviewers' task is to reconstruct the rationale behind each decision based solely on this limited information. The goal is to achieve a high level of agreement (e.g., ⩾90%) with the ground truth, indicating that the system's actions are consistently interpretable without recourse to raw data . This process serves as a crucial test of the system's commitment to explainability and informational self-determination, key tenets of neurorights [14]. Feedback from these audits will be used to refine the fields captured in the logs and improve the quality of explanations generated by the system, ensuring that users and auditors can understand the basis for any adaptive intervention .

Beyond initial deployment, the system incorporates a continuous monitoring and incident response framework managed through the CI/CD pipelines and predefined rollback mechanisms . The CI pipelines for both the ALN contracts and the Rust/C++ codebase are configured to act as gatekeepers, automatically rejecting any changes that weaken existing invariants or introduce coercive controls—for example, by attempting to remove an opt-out flag from the rights contract . Furthermore, every change must be accompanied by proofs (i.e., passing tests and cryptographic stamps) demonstrating that it preserves or tightens safety and eco bounds . To handle unforeseen issues, the system includes predefined ALN rollback shards. These special shards can be deployed to globally or per-cohort disable specific adaptations if monitoring detects a regression or potential harm, always ensuring that the user's right to opt-out remains intact . This combination of automated CI validation, empirical pilot data, transparent auditing, and rapid rollback capabilities creates a closed-loop system for continuous improvement and risk mitigation, ensuring that the bioscale safety system remains trustworthy and effective throughout its operational lifespan.

Strategic Synthesis and Future Directions

The analysis confirms that the proposed non-implant bioscale safety system represents a highly sophisticated and novel approach to regulating AI systems, particularly in the sensitive domain of neurotechnology. Its core strength lies in the tight coupling of governance and engineering, where high-level principles are translated into machine-readable contracts and subsequently enforced through a provably safe toolchain built on Rust and C++. This dual-track model successfully bridges the gap between abstract ethical ideals and concrete, verifiable software artifacts. The system is not merely compliant with regulations; it is designed to be *provably* safe and rights-respecting by construction. This is evident in its technical enforcement of key invariants: the use of bounded types in Rust to prevent invalid data, the one-way aggregation in C++ to protect mental privacy, and the monotonicity constraints to ensure that system improvements do not come at the cost of safety or environmental well-being .

The system's architecture demonstrates a mature understanding of global regulatory landscapes. It is deliberately aligned with the EU AI Act's stringent requirements for high-risk systems, covering risk management, data governance, human oversight, and societal impact [35] [37]. It preemptively satisfies the core tenets of emerging neurorights laws, such as Chile's protections for mental privacy, cognitive liberty, and mental integrity [2]. Furthermore, its design philosophy resonates with the guiding principles of the OECD and UNESCO, emphasizing responsible innovation, anticipatory governance, and the safeguarding of the human mind [19] [50]. This multi-faceted alignment significantly de-risks the system for deployment across different jurisdictions and builds a strong case for its ethical legitimacy.

However, while the technical and governance architecture is robust, the provided materials highlight several critical areas that require deeper investigation and validation before the system can be considered fully mature. The most prominent gap is in the methodology for practical bias auditing. While the plan to analyze break rates and throttle decisions for fairness is a necessary step, the methods for defining demographic variables, controlling for correlated attributes, and interpreting results from de-identified logs are not detailed . This remains a complex challenge, as even aggregated data can reveal biases that are difficult to detect and mitigate [7] [52]. Second, the initial human trials, such as the Phoenix pilot, are necessarily limited in scope, focusing on a small cohort over a short duration . Long-term studies are essential to understand potential habituation effects, shifts in user behavior, and any unforeseen psychological impacts that might arise from prolonged use. The link between reduced device hours and actual, sustained improvements in user well-being needs more robust measurement than self-reporting alone can provide. Third, the legal and liability framework for the system, especially concerning jurisdictional differences, is acknowledged but not fully fleshed out. The `jurisdictiontags` field in the ALN shard is a nod to this complexity, but a clear legal strategy addressing liability in cases of system failure is needed to complement the technical safeguards [2] [4]. Finally, the security of the local processing environment on commodity devices, while mentioned, warrants deeper scrutiny. Protecting against vulnerabilities like data poisoning during the training phase is a known concern that requires dedicated mitigation strategies [7].

In conclusion, the bioscale safety system presents a compelling blueprint for the future of regulated AI. Its innovative use of ALN contracts and a provably safe toolchain offers a powerful model for building trustworthy neurotechnologies. The system successfully translates complex governance requirements into implementable technical specifications. The path forward involves moving from this robust theoretical foundation to rigorous, long-term empirical validation. Addressing the identified gaps in bias auditing methodologies, conducting longitudinal human studies, clarifying the legal and liability framework, and performing deep security assessments will be paramount. Successfully navigating these challenges will solidify the system's position not just as a technological achievement, but as a responsible and ethical paradigm for human-computer interaction in the age of AI.

Reference

1. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
2. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
3. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
4. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
5. Arxiv今日论文| 2025-12-11 http://lonepatient.top/2025/12/11/arxiv_papers_2025-12-11
6. Computer Vision and Pattern Recognition May 2025 https://www.arxiv.org/list/cs.CV/2025-05?skip=125&show=2000
7. 人工智能2025_10_30 https://www.arxivdaily.com/thread/73297

8. COLING 2025 The 31st International Conference on ... https://aclanthology.org/2025.coling-industry.pdf
9. HUAWEI RESEARCH https://www-file.huawei.com/admin/asset/v1/pro/view/a4999ca41b584ec1b9da65fab98d87d6.pdf
10. Implications of the novel EU AI Act for neurotechnologies https://www.sciencedirect.com/science/article/pii/S089662732400607X
11. Assessing High-Risk AI Systems under the EU AI Act https://arxiv.org/html/2512.13907v2
12. The Academic Guide to AI Act Compliance https://hal.science/hal-05365570v1/file/The%20Academic%20Guide%20to%20AI%20Act%20Compliance%20%20-%202025%20-%20Ed.%20MHODAC%20%26%20CP.pdf
13. TechDispatch #1/2024 - Neurodata https://www.edps.europa.eu/data-protection/our-work/publications/techdispatch/2024-06-03-techdispatch-12024-neurodata
14. Neurorights, Mental Privacy, and Mind Reading | Neuroethics https://link.springer.com/article/10.1007/s12152-024-09568-z
15. Neurotechnology Toolkit https://www.oecd.org/content/dam/oecd/en/topics/policy-sub-issues/emerging-technologies/neurotech-toolkit.pdf
16. Framework for Anticipatory Governance of Emerging ... https://www.oecd.org/content/dam/oecd/en/publications/reports/2024/04/framework-for-anticipatory-governance-of-emerging-technologies_14bf0402/0248ead5-en.pdf
17. Responsible innovation in neurotechnology enterprises https://www.oecd.org/en/publications/responsible-innovation-in-neurotechnology-enterprises_9685e4fd-en.html
18. Neuroethics guidance documents: principles, analysis, and ... https://pmc.ncbi.nlm.nih.gov/articles/PMC10602660/
19. Ethics of neurotechnology: UNESCO adopts the first global ... https://www.unesco.org/en/articles/ethics-neurotechnology-unesco-adopts-first-global-standard-cutting-edge-technology
20. Draft text of the Recommendation on the Ethics ... https://unesdoc.unesco.org/ark:/48223/pf0000393395
21. First draft of the Recommendation on the Ethics ... https://unesdoc.unesco.org/ark:/48223/pf0000391444
22. Ethics of neurotechnology https://www.unesco.org/en/ethics-neurotech
23. Final report on the draft Recommendation on the Ethics of ... https://unesdoc.unesco.org/ark:/48223/pf0000393266
24. unesco https://research.bjmu.edu.cn/docs/2024-10/af391e2749cf4794b652945bebbb5b91.pdf

25. Draft Recommendation on the Ethics of Neurotechnology https://unesdoc.unesco.org/ark:/48223/pf0000394866
26. UNESCO adopts ethics framework for neurotechnology https://www.linkedin.com/posts/bambi1_ethicalai-neurotechnology-unesco-activity-7394682114161639424-CbwB
27. A/80/283 - General Assembly - the United Nations https://docs.un.org/en/A/80/283
28. First draft of a Recommendation on the Ethics ... https://unesdoc.unesco.org/ark:/48223/pf0000391074
29. Arxiv今日论文 | 2025-12-16 | 闲记算法 http://lonepatient.top/2025/12/16/arxiv_papers_2025-12-16
30. A Systematic Literature Review on Detecting Software ... https://arxiv.org/html/2507.22659v1
31. Findings of the Association for Computational Linguistics https://aclanthology.org/volumes/2025.findings-acl/
32. Arxiv今日论文 | 2025-11-03 http://lonepatient.top/2025/11/03/arxiv_papers_2025-11-03
33. Artificial Intelligence https://arxiv.org/list/cs.AI/new
34. 自然语言处理2025_4_22 http://arxivdaily.com/thread/66586
35. Assessing High-Risk Systems: An EU AI Act Verification ... https://arxiv.org/html/2512.13907v1
36. Harmonised Standards for the European AI Act https://publications.jrc.ec.europa.eu/repository/bitstream/JRC139430/JRC139430_01.pdf
37. The EU AI Act Compliance Playbook: A Step-by- ... https://www.linkedin.com/pulse/eu-ai-act-compliance-playbook-step-by-step-guide-high-risk-khan-g7htf
38. Recommendations for Responsible Development and ... https://pmc.ncbi.nlm.nih.gov/articles/PMC8081770/
39. Neuroethics guidance documents: principles, analysis, and ... https://academic.oup.com/jlb/article/10/2/lsad025/7329870
40. Neurotechnology https://www.ohchr.org/sites/default/files/documents/hrbodies/hrcouncil/advisorycommittee/neurotechnology/03-ngos/ac-submission-cso-neurorightsfoundation.pdf
41. Computer Safety, Reliability, and Security https://link.springer.com/content/pdf/10.1007/978-3-030-55583-2.pdf
42. Communication and Intelligent Systems https://www.researchgate.net/profile/Vamsi-Nagaraju/publication/352811793_Prediction_of_California_Bearing_Ratio_of_Subgrade_Soils_Using_Artificial_Neural_Network_Principles/links/634fe41896e83c26eb357781/Prediction-of-

California-Bearing-Ratio-of-Subgrade-Soils-Using-Artificial-Neural-Network-Principles.pdf

43. Pccda Book | PDF | Artificial Intelligence https://www.scribd.com/document/790683919/Pccda-Book
44. Draft Recommendation on the Ethics of Neurotechnology https://unesdoc.unesco.org/ark:/48223/pf0000394861
45. UNESCO Adopts First Global Framework on ... https://www.globalpolicywatch.com/2026/01/unesco-adopts-first-global-framework-on-neurotechnology-ethics/
46. Towards an International Instrument https://www.unesco.org/en/ethics-neurotech/recommendation
47. Engineering Trustworthy Machine-Learning Operations ... https://www.arxiv.org/pdf/2505.20136
48. Final report on the draft text of the Recommendation ... https://unesdoc.unesco.org/ark:/48223/pf0000393400
49. 20061030 | PDF https://www.scribd.com/doc/65946161/20061030
50. OECD Science, Technology and Innovation Outlook 2025 https://www.oecd.org/en/publications/oecd-science-technology-and-innovation-outlook-2025_5fe57b90-en/full-report/technology-convergence-trends-prospects-and-policies_5902a07e.html
51. responsible business conduct and anticipatory governance ... https://www.oecd.org/content/dam/oecd/en/publications/reports/2025/06/responsible-business-conduct-and-anticipatory-governance-of-emerging-technology_ebc54918/1308a723-en.pdf
52. Regulating for the future: OECD Regulatory Policy Outlook ... https://www.oecd.org/en/publications/oecd-regulatory-policy-outlook-2025_56b60e39-en/full-report/regulating-for-the-future_e948d334.html
53. OECD Science, Technology and Innovation Outlook 2021 https://www.oecd.org/content/dam/oecd/en/publications/reports/2021/01/oecd-science-technology-and-innovation-outlook-2021_3f424d14/75f79015-en.pdf
54. agile mechanisms for responsible technology development https://www.oecd.org/content/dam/oecd/en/publications/reports/2025/04/agile-mechanisms-for-responsible-technology-development_84f0fcea/2a35358e-en.pdf
55. Integrity and security in the global research ecosystem (EN) https://www.oecd.org/content/dam/oecd/en/publications/reports/2022/06/integrity-and-security-in-the-global-research-ecosystem_2bd8511d/1c416f43-en.pdf
56. International Journal of Learning, Teaching And ... https://www.researchgate.net/profile/Fatima-Belkhir/publication/377151608_Vol_22_No_8_August_2023/links/6596f3580bb2c7472b32c5fe/Vol-22-No-8-August-2023.pdf

57. Dying To Be Ill - True Stories of Medical Deception - Marc D https://www.scribd.com/document/810738207/Dying-to-be-Ill-True-Stories-of-Medical-Deception-Marc-D-Feldman-Gregory-P-Yates-2018-Taylor-and-Francis-9781138063815-edaafd2f9644
58. China https://www.nsfc.gov.cn/Portals/0/fj/fj20230220_01.xlsx
59. 978-3-031-57246-3 | PDF | Formal Verification | Computing https://www.scribd.com/document/977584758/978-3-031-57246-3
60. Big Data Analytics in Astronomy, Science, and Engineering https://link.springer.com/content/pdf/10.1007/978-3-031-86193-2.pdf
61. Emerging technologies https://www.oecd.org/en/topics/emerging-technologies.html
62. Shaping a rights-oriented digital transformation (EN) https://www.oecd.org/content/dam/oecd/en/publications/reports/2024/06/shaping-a-rights-oriented-digital-transformation_30378a18/86ee84e2-en.pdf
63. Deliverable 1.2 AI-Medical Device software compliance ... https://ec.europa.eu/research/participants/documents/downloadPublic?documentIds=080166e5e47b3206&appId=PPGMS
64. Navigating the EU AI Act: A Methodological Approach ... https://arxiv.org/html/2403.16808v2
65. Navigating artificial general intelligence development https://www.nature.com/articles/s41598-025-92190-7
66. 43 C/5, Volume 1: Draft resolutions: first biennium of the ... https://unesdoc.unesco.org/ark:/48223/pf0000392982
67. UNESCO's reaches final stages in building a global ... https://www.unesco.org/en/articles/unescos-reaches-final-stages-building-global-standard-ethics-neurotechnology
68. Compilation of comments received from Member States on ... https://unesdoc.unesco.org/ark:/48223/pf0000393396.locale=en